

# Geometric deep learning for data on manifolds, and spherical images

Half-way seminar,

Oscar Carlsson, Department of Mathematical Sciences

For the Geometric Deep Learning seminar series on 2023-02-09,  
at the Department of Mathematical Sciences

**WASP**

WALLENBERG AI,  
AUTONOMOUS SYSTEMS  
AND SOFTWARE PROGRAM



**CHALMERS**  
UNIVERSITY OF TECHNOLOGY



UNIVERSITY OF GOTHEBURG

# Material

This presentation is based on the following papers:

- “**Geometric Deep Learning and Equivariant Neural Networks**” by Jan E. Gerken, Jimmy Aronsson, Oscar Carlsson, Hampus Linander, Fredrik Ohlsson, Christoffer Petersson, Daniel Persson. *Under review for AIRE*.
- “**Equivariance versus Augmentation for Spherical Images**” by Jan E. Gerken, Oscar Carlsson, Hampus Linander, Fredrik Ohlsson, Christoffer Petersson, Daniel Persson. *Accepted for ICML22*.

# Outline

- 1 General machine learning introduction and background
- 2 Introduction to geometric deep learning
- 3 Papers:
  - 1 Paper I: “Geometric Deep Learning and Equivariant Neural Networks”
    - 1 Overview
    - 2 Mathematical background
    - 3 Gauge equivariant convolution
  - 2 Paper II: “Equivariance versus Augmentation for Spherical Images”
    - 1 Background
    - 2 Results
    - 3 Conclusions and takeaways
- 4 Future work

# Introduction

- Machine learning

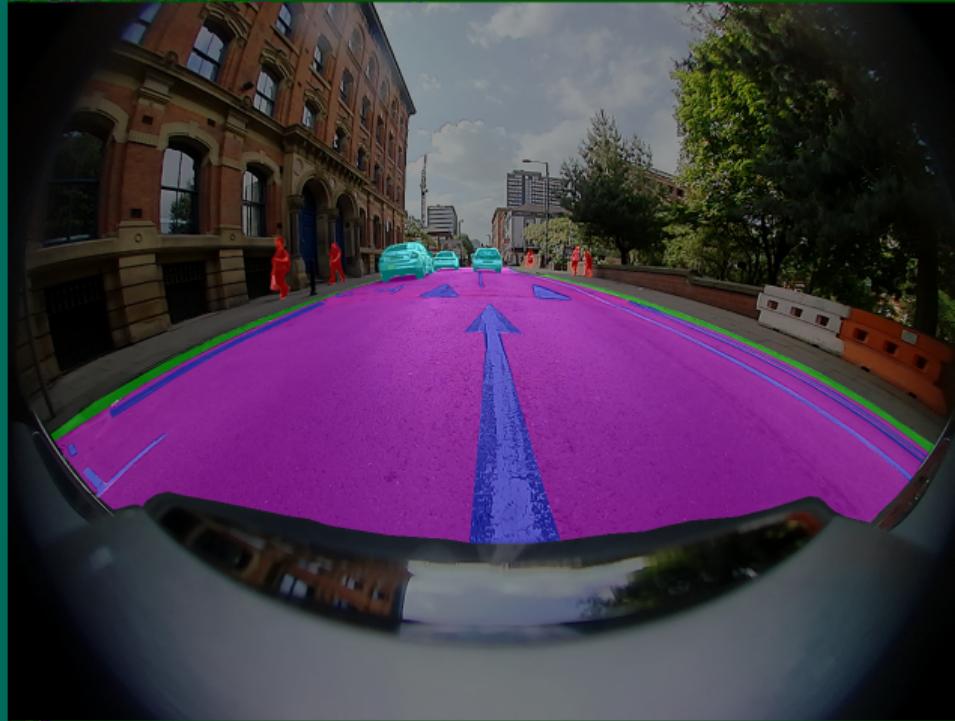


# Examples of machine learning applications

- Language models
- AlphaFold
- Generative models
- Image analysis
  - Object detection
  - Semantic segmentation
  - Depth estimation
  - Classification



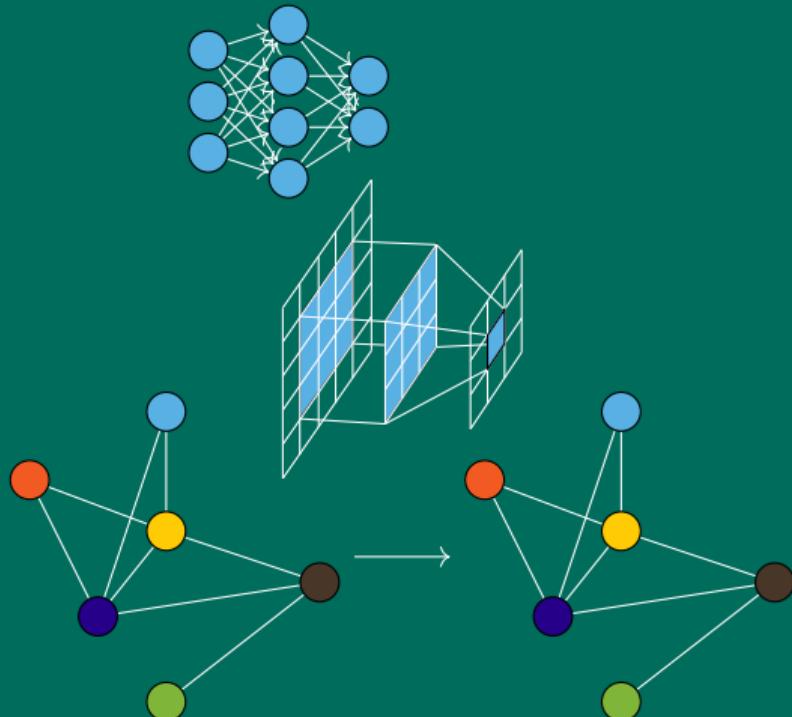
# Image analysis example



From the WoodScape dataset (Yogamani et al. 2019).

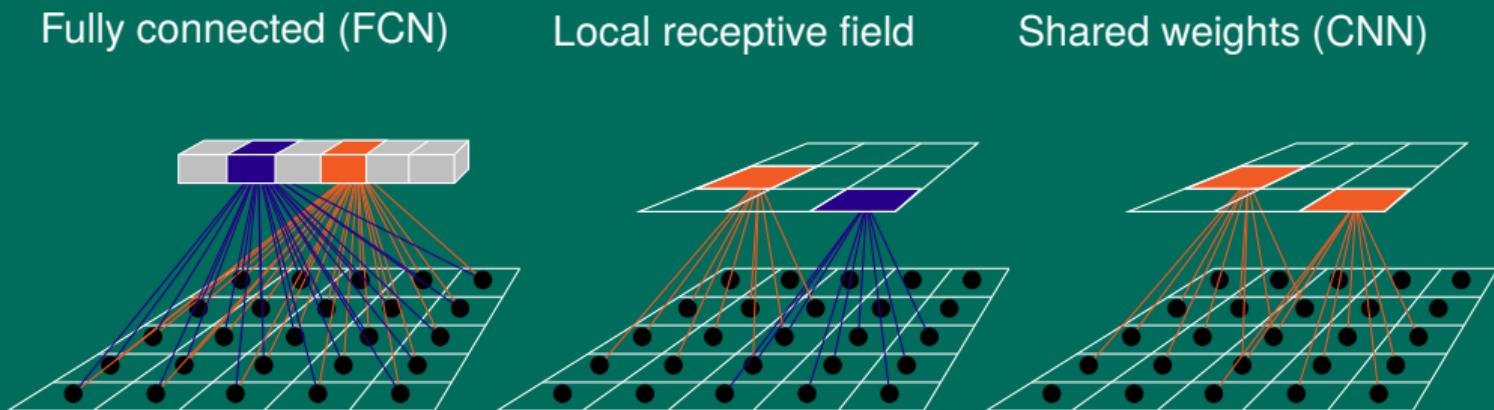
# Some architecture examples

- Fully connected networks (FCN)
- Convolutional neural networks (CNN)
- Graph networks

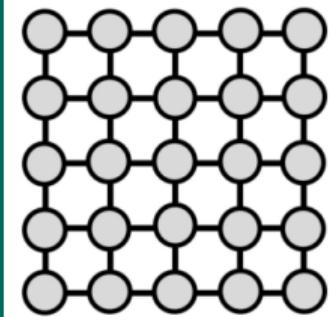


# FCN vs CNN

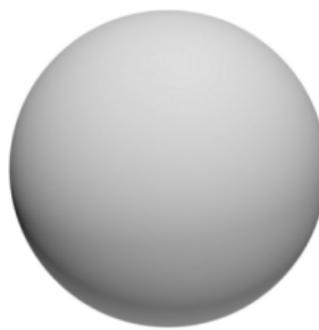
Q: Why does FCN not work as well as CNN for image analysis?



# What is geometric deep learning?



**Grids**



**Groups**



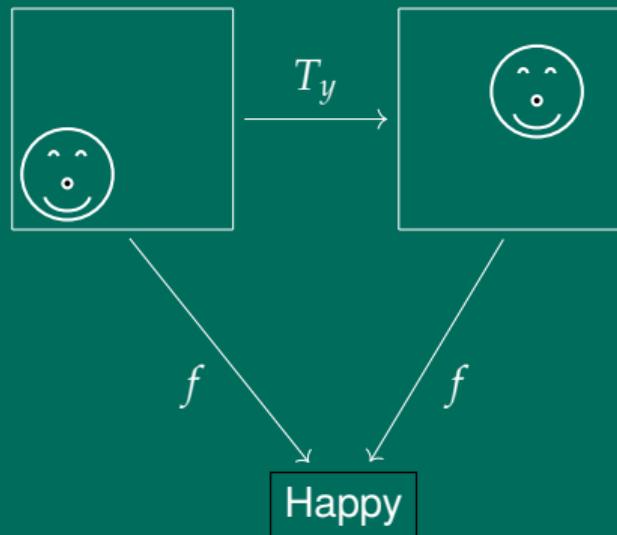
**Graphs**



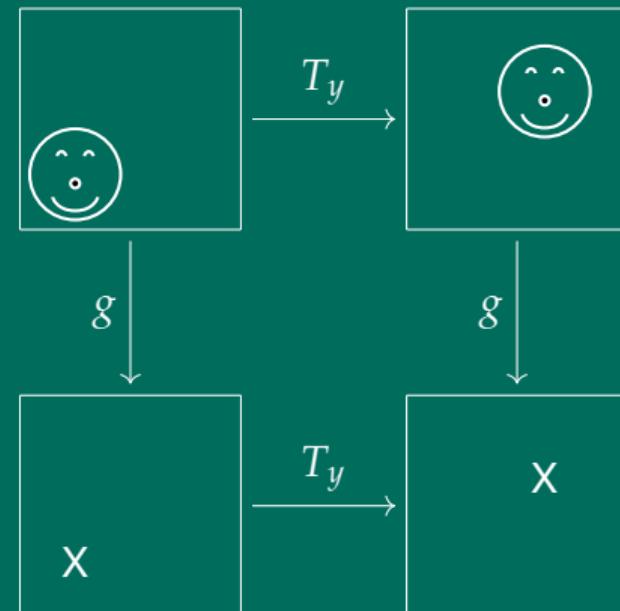
**Geodesics & Gauges**

(Bronstein et al. 2021)

# Invariance and equivariance



$$(T_yf)(x) = f(x)$$



$$(T_yg)(x) = g(T_yx)$$

# Group actions and representations

**Important viewpoint/assumption:** data is some function on an underlying space. In this viewpoint *layers maps data functions to data functions*.

E.g. Viewing input colour image as a function  $f : \mathbb{R}^2 \rightarrow \mathbb{R}^3$ .

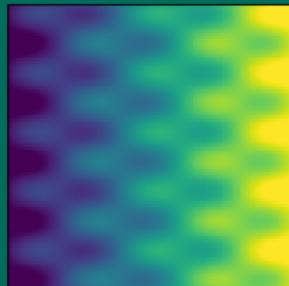
Group of transformations  $G$  acting on the input data, e.g.

$$(g \triangleright f)(x) = \rho(g)f(g^{-1} \triangleright x)$$

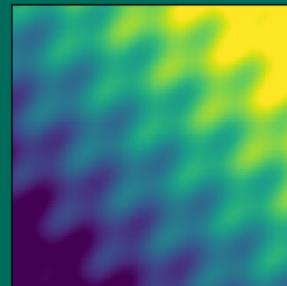
How the data function  $f$  transforms is given by  $\rho$ , sometimes this is referred to as different *data types*.

## Example of data types: scalar

$$\phi(x)$$

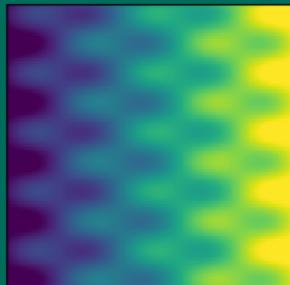


$$(L_g \phi)(x) = \phi(g^{-1}x)$$

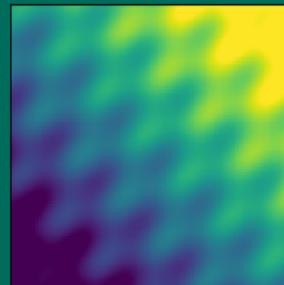


# Example of data types: scalar and vector

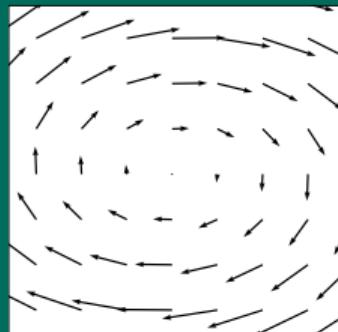
$$\phi(x)$$



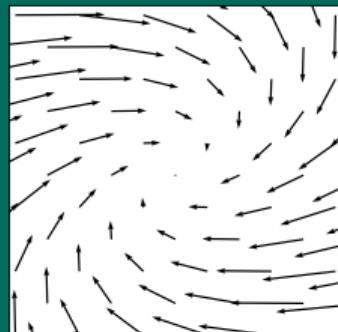
$$(L_g \phi)(x) = \phi(g^{-1}x)$$



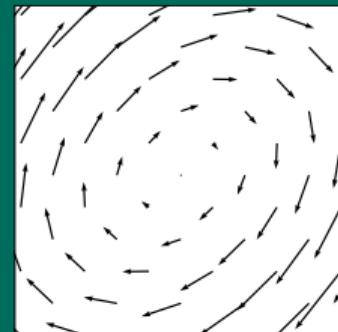
$$f(x)$$



$$f(g^{-1}x)$$

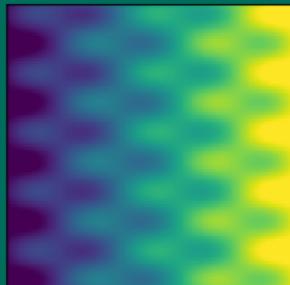


$$(L_g f)(x)$$

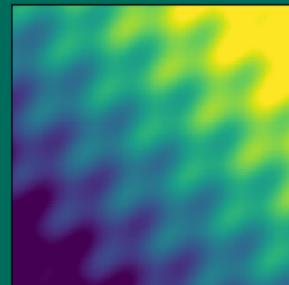


# Example of data types: scalar and vector

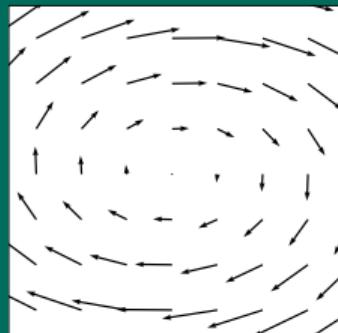
$$\phi(x)$$



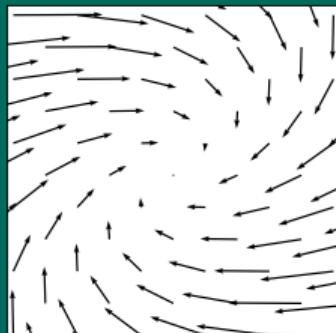
$$(L_g \phi)(x) = \phi(g^{-1}x)$$



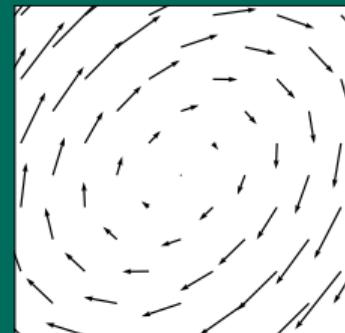
$$f(x)$$



$$f(g^{-1}x)$$

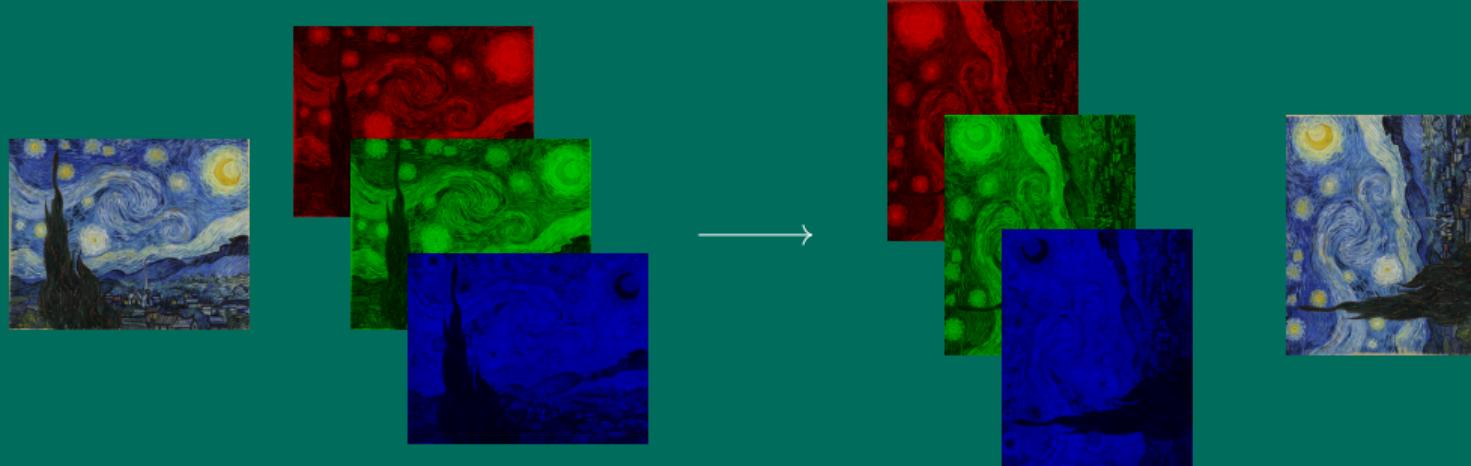


$$(L_g f)(x) = \rho(g)f(g^{-1}x)$$



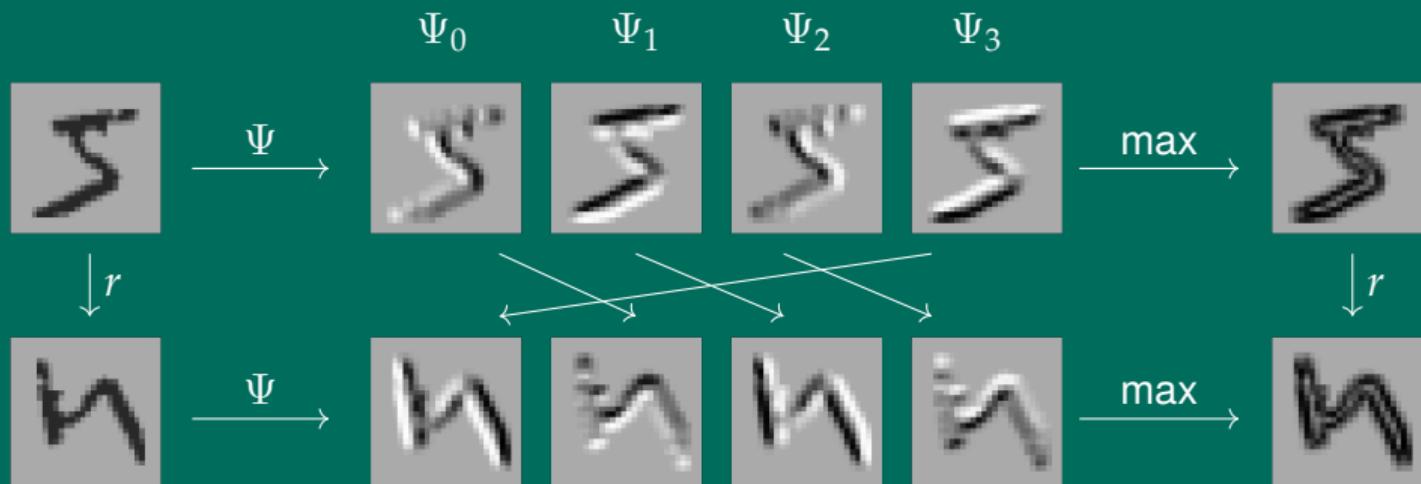
# Example in practice: ordinary image

Group  $G = \text{SO}(2)$  acting as  $(g \triangleright f)(x) = \rho(g)f(g^{-1} \triangleright x)$  where  $\rho = \bigoplus_{i=1}^3 \text{id}$  where  $\text{id}$  is the trivial representation.



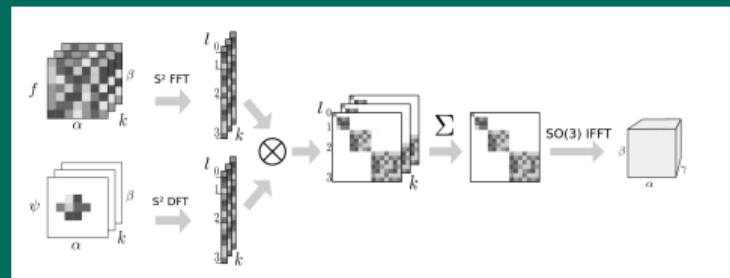
# Regular representation approach

- Copy and rotate filters (Bekkers et al. 2018; T. S. Cohen and Welling 2016)



# Continuous groups: spectral methods

- Spherical cnn: S2CNN (T. S. Cohen, Geiger et al. 2018)



$$[\psi \star f](R) = \langle L_R \psi, f \rangle = \int_{S^2} \sum_{k=1}^K \psi_k(R^{-1}x) f_k(x) dx, \quad R \in SO(3)$$

# Solving the kernel constraint

- Kernel constraint for steerable kernels  $k : \mathbb{R}^2 \rightarrow \mathbb{R}^{\text{dim}_{\text{out}} \times \text{dim}_{\text{in}}}$

$$k(gx) = \rho_{\text{out}}(g)k(x)\rho_{\text{in}}(g^{-1}), \quad \forall g \in G$$

change basis to irreps and redefine kernel  $k$  gives

$$k(gx) = \left( \bigoplus_{i \in I_{\text{out}}} \psi_i(g) \right) k(x) \left( \bigoplus_{j \in I_{\text{in}}} \psi_j^{-1}(g) \right).$$

(Weiler and Cesa 2019)

# The papers

- Paper I: Geometric Deep Learning and Equivariant Neural Networks  
(Gerken, Aronsson et al. 2021)
- Paper II: Equivariance versus Augmentation for Spherical Images  
(Gerken, Carlsson et al. 2022)

# Paper I: Geometric Deep Learning and Equivariant Neural Networks

## -Overview

Covers mathematical approaches to equivariant networks for data defined on a continuous (possibly discretised) spaces covering

- General (Riemannian) manifolds
- Meshes
- General Homogenous spaces
- Spectral networks on the sphere

# Gauge equivariant convolutions on Riemannian manifolds

Lifting the framework introduced in T. Cohen et al. 2020 and refined in Aronsson 2021 to a general Riemannian manifold where the transformations are local coordinate changes.

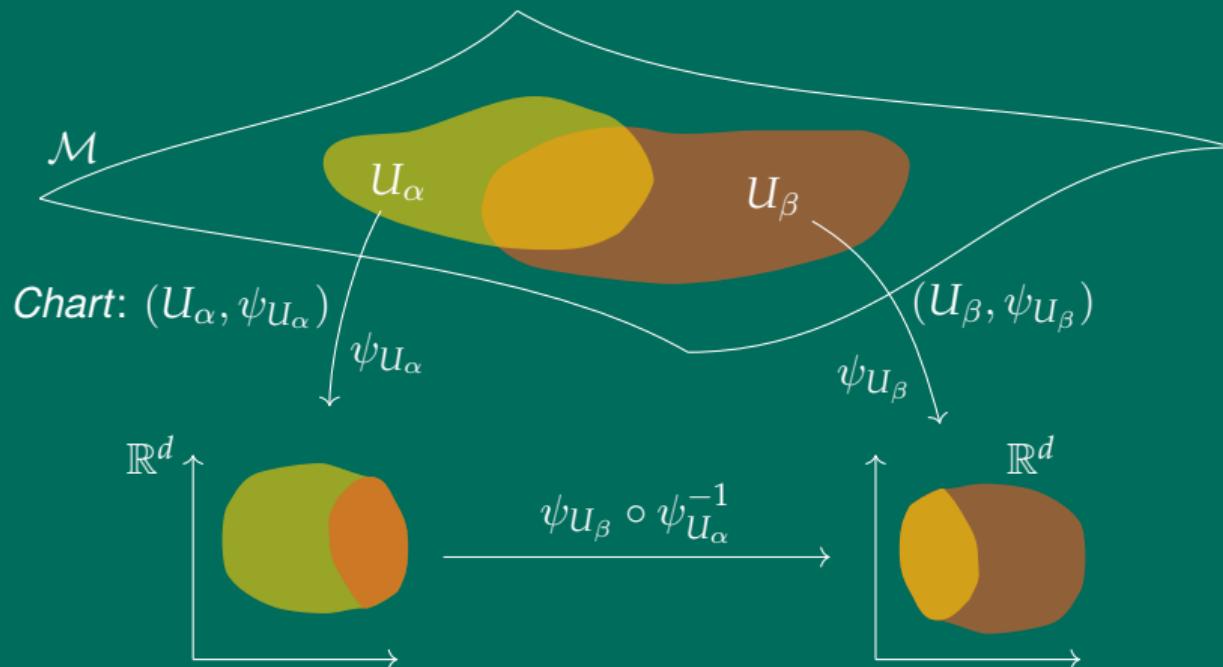
Necessary concepts:

- Manifolds
- Fibre bundles (vector bundles, principal bundles, and associated bundles)
- Gauge transformations
- Parallel transport

Gauge equivariant convolution  $\phi : C_c(P, \rho) \rightarrow C_c(P, \eta)$  by

$$(\phi f)(p) = [\kappa \star f](p) = \int_{B_R} \kappa(X) f(\mathcal{T}_X p) \text{vol}_{T_x \mathcal{M}} .$$

# Recap on manifolds

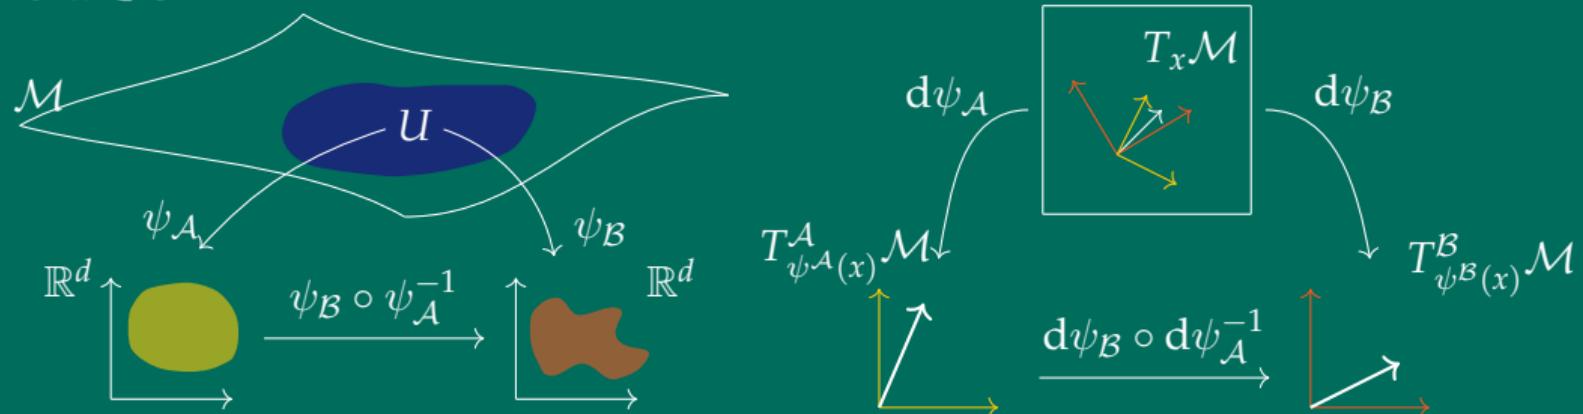


Atlas:  $\{(U_\alpha, \psi_{U_\alpha}) : \alpha \in I\}$ , transition maps  $\tau_{\alpha,\beta} = \psi_{U_\beta} \circ \psi_{U_\alpha}^{-1}$  for  $U_\alpha \cap U_\beta \neq \emptyset$

# Gauge transformations (change of local coordinates)

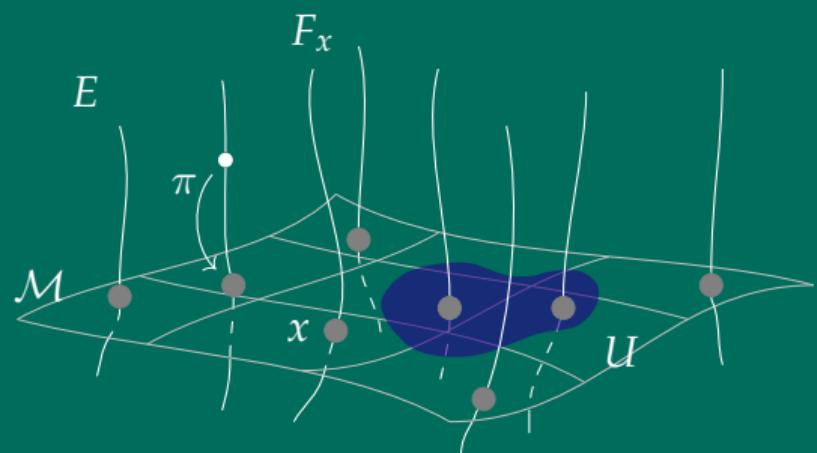
A chart  $(U, \psi_A)$  is an assignment of local coordinates to  $U \subseteq \mathcal{M}$ , choosing another chart  $(U, \psi_B)$  gives a different set of local coordinates.

The map  $\psi_B \circ \psi_A^{-1}$  induces a gauge transformation in the tangent spaces  $T_x \mathcal{M} \cong \mathbb{R}^d$  for  $x \in U$ .



# Recap on bundles

- *Bundle*: A bundle  $(E, \pi, \mathcal{M})$ ,  $E \xrightarrow{\pi} \mathcal{M}$ , is a triple consisting of a (total) space  $E$ , a surjective continuous map  $\pi : E \rightarrow \mathcal{M}$  (projection), and a base manifold  $\mathcal{M}$ .
- *Vector bundle*: Each fibre is a vector space  $V$ 
  - Ex: Tangent bundle  $T\mathcal{M}$ ,  $T_x\mathcal{M} \cong \mathbb{R}^d$ ,  $\forall x \in \mathcal{M}$
- *Principal  $K$ -bundle*: Each fibre is homeomorphic to the group  $K$



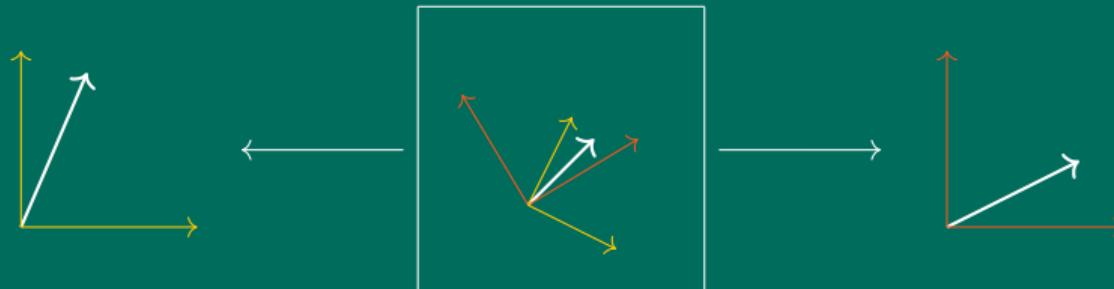
# Associated vector bundles

*Associated bundle:* A bundle constructed from a principal  $K$ -bundle  $P$  and a representation  $(V, \rho)$ , i.e.  $k \triangleright v = \rho(k)v$ , through the equivalence relation

$$(p, v) \sim_{\rho} (p \triangleleft k, k^{-1} \triangleright v)$$

on the space  $P \times V$ . The quotient  $P \times V / \sim_{\rho} =: E_{\rho}$  with the projection  $\pi([p, v]) = \pi_P(p)$  makes  $E_{\rho}$  to a vector bundle over the same base space as  $P$ .

**Intuition:**  $[p, v]$  is a geometric (i.e. coordinate free) vector. A specific instance  $(p, v) \in [p, v]$  expresses the geometric vector in a “frame”/basis  $p$  and components  $v$ .



# Framework used in Aronsson 2021

- *Data point*: section  $\sigma \in \Gamma(E_\rho)$  of associated bundle
- *Feature map*: map  $f \in C_c(P, \rho)$  satisfying

$$k \triangleright f(p) = f(p \triangleleft k^{-1})$$

or  $k \circ f = f \circ k^{-1}$ .

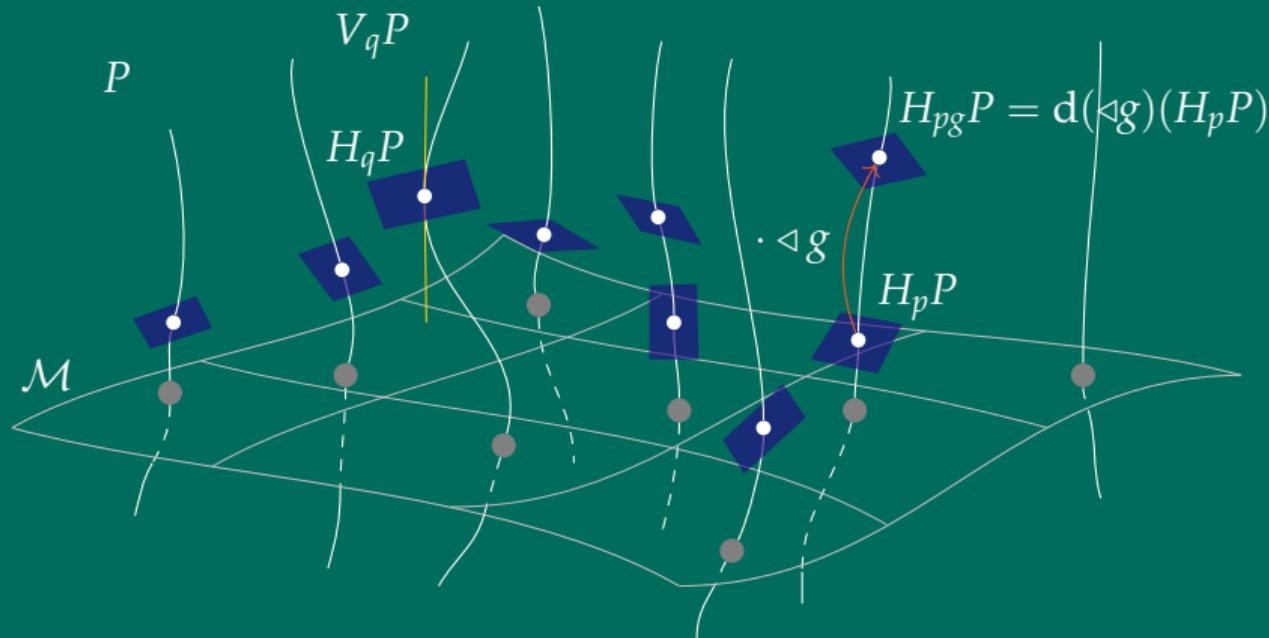
- **Theorem**  $\varphi_\rho : C_c(P, \rho) \rightarrow \Gamma(E_\rho)$  acting by

$$f \mapsto \sigma_f = [\pi^{-1}, f \circ \pi^{-1}]$$

is an isomorphism between  $C_c(P, \rho)$  and  $\Gamma(E_\rho)$  (Kolar et al. 1993).

# Connection and parallel transport

**Principal Ehresmann connection:** a smooth subbundle  $H$  of  $TP$  such that  $H \oplus V = TP$  where  $V = \ker(d\pi : TP \rightarrow T\mathcal{M})$ . At each  $p \in P$  then  $T_p P = H_p \oplus V_p$



# Connection and parallel transport

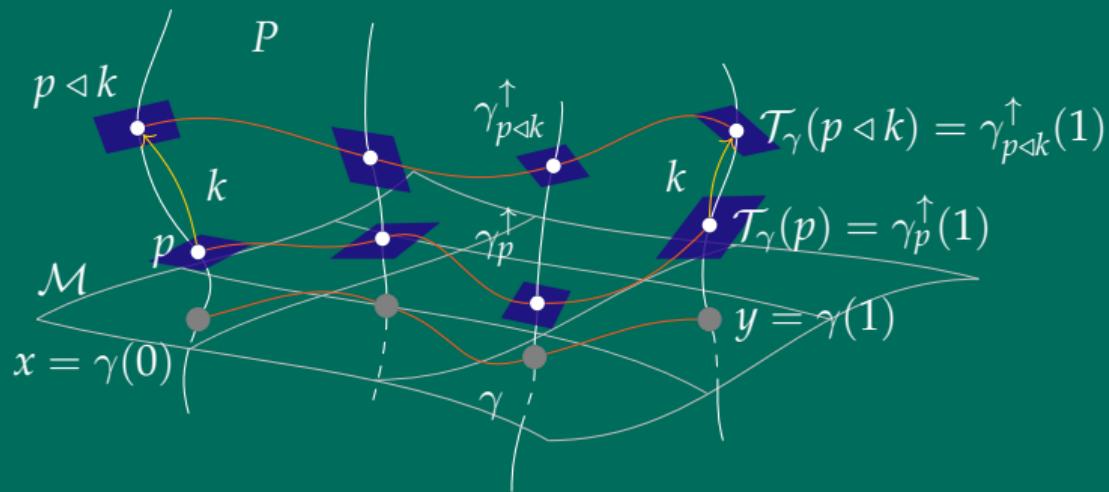
Horizontal lift  $\gamma_p^\uparrow : [0, 1] \rightarrow P$  of curve  $\gamma : [0, 1] \rightarrow \mathcal{M}$  s.t.  $\gamma(0) = x$  and  $\gamma(1) = y$ .

Parallel transport in  $P$  by

$$\mathcal{T}_\gamma p = \gamma_p^\uparrow(1),$$

which is equivariant to the principal action

$$\mathcal{T}_\gamma(p) \triangleleft k = \mathcal{T}_\gamma(p \triangleleft k).$$



# Gauge equivariant convolution

Kernel:  $\kappa : T\mathcal{M} \rightarrow \text{Hom}(E_\rho, E_\eta)$ .

In local coordinates  $\mathcal{A}$  we have

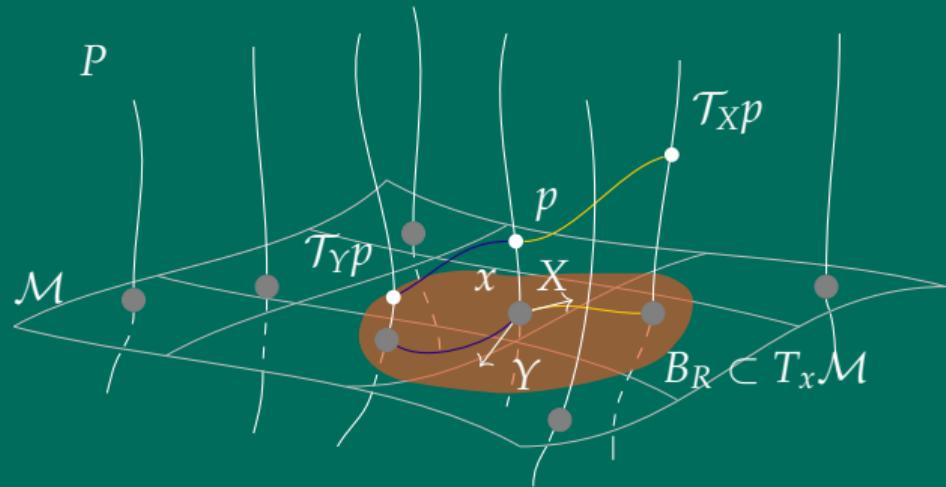
$$\kappa^{\mathcal{A}} : \mathbb{R}^d \rightarrow \mathbb{R}^{\dim(\eta) \times \dim(\rho)}$$

transforming as

$$\kappa^{\mathcal{A}}(X^{\mathcal{A}}) = \eta(k^{-1})\kappa^{\mathcal{B}}(X^{\mathcal{B}})\rho(k)$$

where  $X^{\mathcal{B}} = k \triangleright X^{\mathcal{A}}$ . Then  $\phi :$   
 $C_c(P, \rho) \rightarrow C_c(P, \eta)$  by

$$(\phi f)(p) = [\kappa \star f](p) = \int_{B_R} \kappa(X)f(\mathcal{T}_X p) \text{vol}_{T_x \mathcal{M}}.$$



# Calculations

Evaluate the convolution at a frame  $p^{\mathcal{A}}$  and let  $k \in K$  be s.t.  $p^{\mathcal{B}} = p^{\mathcal{A}} \triangleleft k^{-1}$  then

$$\begin{aligned} k \triangleright (\phi f)(p^{\mathcal{A}}) &= k \triangleright [\kappa^{\mathcal{A}} \star f](p^{\mathcal{A}}) = \eta(k) \int_{B_R^{\mathcal{A}}} \kappa^{\mathcal{A}}(X^{\mathcal{A}}) f(\mathcal{T}_X p^{\mathcal{A}}) \text{vol}_{T_x \mathcal{M}}^{\mathcal{A}} \\ &= \int_{B_R^{\mathcal{A}}} \kappa^{\mathcal{B}}(X^{\mathcal{B}}) \rho(k) f(\mathcal{T}_X p^{\mathcal{A}}) \text{vol}_{T_x \mathcal{M}}^{\mathcal{A}} = \left\{ \rho(k) f(p) = f(p \triangleleft k^{-1}) \right\} \\ &= \int_{B_R^{\mathcal{A}}} \kappa^{\mathcal{B}}(X^{\mathcal{B}}) f \left( (\mathcal{T}_X p^{\mathcal{A}}) \triangleleft k^{-1} \right) \text{vol}_{T_x \mathcal{M}}^{\mathcal{A}} = \{ \mathcal{T}_X(p) \triangleleft k = \mathcal{T}_X(p \triangleleft k) \} \\ &= \int_{B_R^{\mathcal{A}}} \kappa^{\mathcal{B}}(X^{\mathcal{B}}) f \left( \mathcal{T}_X(p^{\mathcal{A}} \triangleleft k^{-1}) \right) \text{vol}_{T_x \mathcal{M}}^{\mathcal{A}} \\ &= \int_{B_R^{\mathcal{B}}} \kappa^{\mathcal{B}}(X^{\mathcal{B}}) f(\mathcal{T}_X p^{\mathcal{B}}) \text{vol}_{T_x \mathcal{M}}^{\mathcal{B}} \\ &= (\phi f)(p^{\mathcal{B}}) = (\phi f)(p^{\mathcal{A}} \triangleleft k^{-1}) \quad \Rightarrow \quad \phi f \text{ is a feature map} \end{aligned}$$

$$k \circ \phi \circ f = \phi \circ f \circ k^{-1} = \phi \circ k \circ f \quad \Rightarrow \quad k \circ \phi = \phi \circ k \quad \phi \text{ equivariant}$$

# Some details on the kernel

Defining the kernel as

$$\kappa : T\mathcal{M} \rightarrow \text{Hom}(E_\rho, E_\eta)$$

allows for different kernels at different  $x \in \mathcal{M}$ , i.e. a *kernel field*.

Noting that in local a local coordinate chart  $\mathcal{A}$  at  $x \in \mathcal{M}$  this kernel becomes

$$\mathcal{K}^{\mathcal{A}} : \mathbb{R}^d \rightarrow \mathbb{R}^{\dim(\eta) \times \dim(\rho)}.$$

To get weight sharing over  $x \in \mathcal{M}$  pick

$$K : \mathbb{R}^d \rightarrow \mathbb{R}^{\dim(\eta) \times \dim(\rho)}$$

and set

$$\mathcal{K}^{\mathcal{X}} = K,$$

for any gauge  $\mathcal{X}$ . This kernel can now be shared over different  $x \in \mathcal{M}$

---

(Weiler, Forré et al. 2021)

# Paper II: Equivariance versus Augmentation for Spherical Images

## -Background

### Data augmentation

- Simple to implement
- "Generating more points on the data manifold"
- No guarantee for "exact" equivariance
- Increases training time

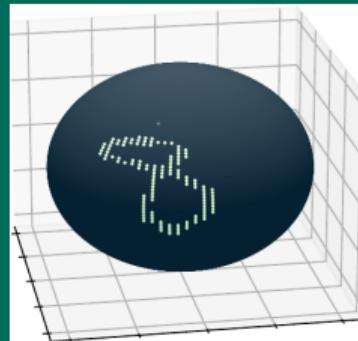
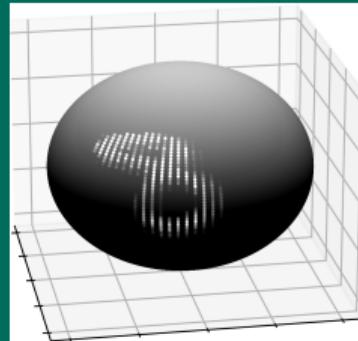
### Equivariance

- Can be more complicated to code with more details in the model
- Theoretically more complicated
- Theoretical insurance that symmetry is respected

**Often claimed:** equivariant networks are better, but little is actually investigated; i.e. what is the limit of *geometrical data augmentation*?

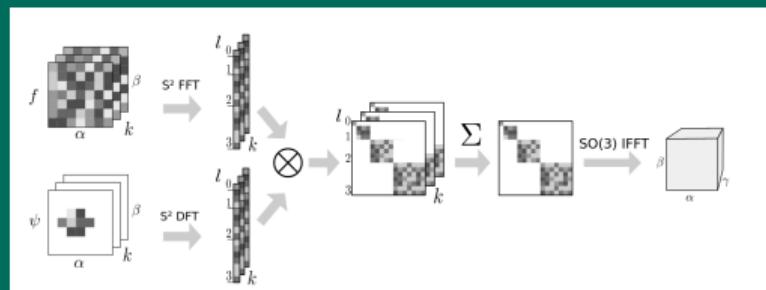
# Paper II: Setup

- Dataset: MNIST images projected to the sphere.
  - Data augmentation: pick random  $g \in \text{SO}(3)$  and act on an image
- Models:
  - S2CNN (T. S. Cohen, Geiger et al. 2018)
  - Ordinary CNN
- Tasks:
  - Classification (invariant)
  - Semantic segmentation (equivariant)

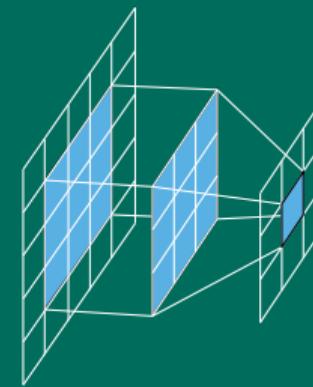


# Paper II: Networks

S2CNN :



CNN:



# Paper II: Network selection

20 random networks for each parameter range and type.

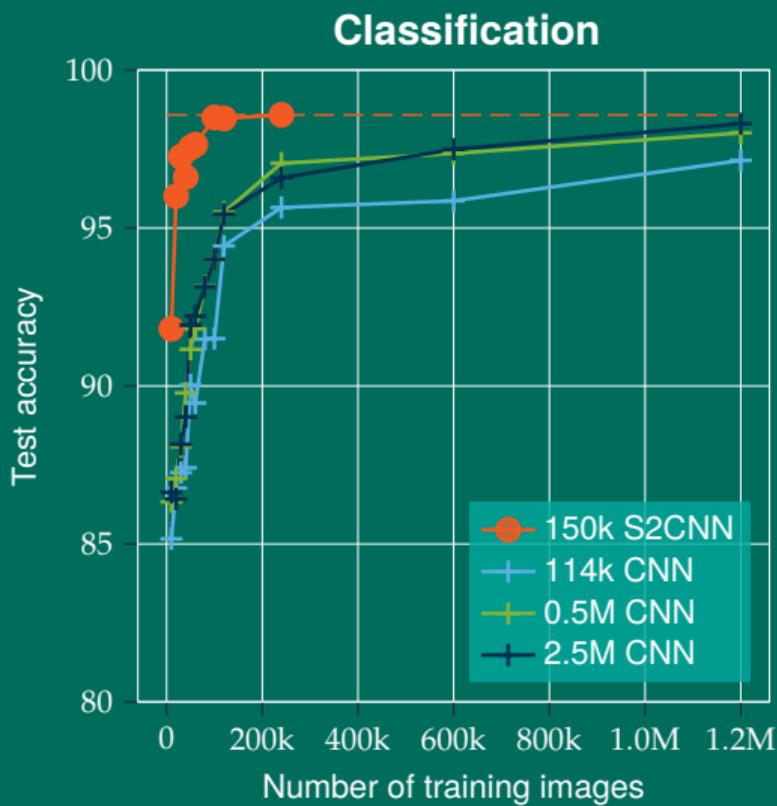
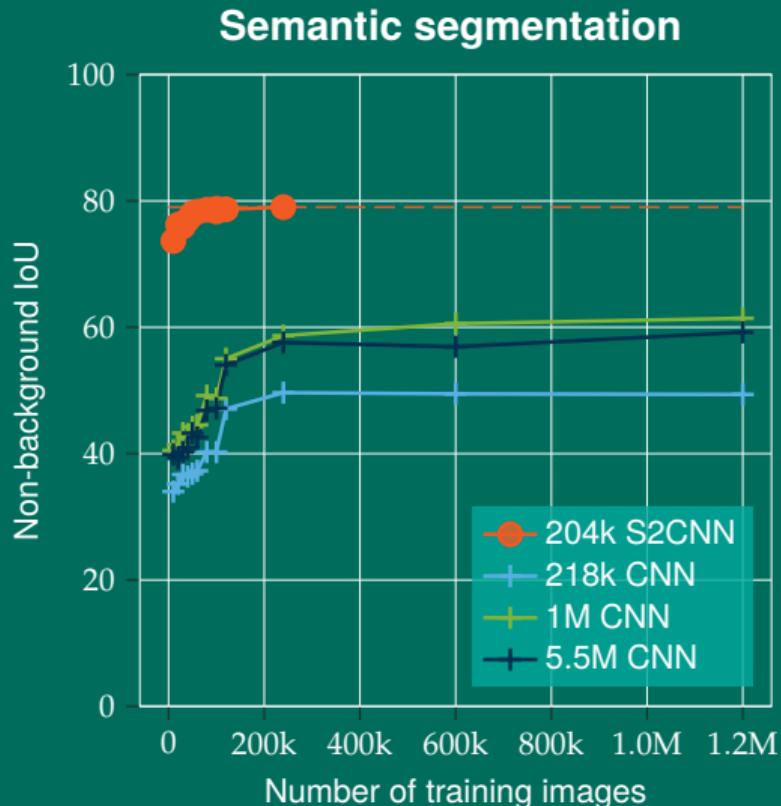
## S2CNN

- Semantic segmentation
  - 204k
  - (820k)
- Classification
  - 150k

## CNN

- Semantic segmentation
  - 218k
  - 1M
  - 5.5M
- Classification
  - 114k
  - 0.5M
  - 2.5M

# Paper II: Some results



# Paper II: Latency and training times

## Latency

Model	Latency (ms)	Throughput (N/s)
204k S2CNN	$111 \pm 0.6$	$9.0 \pm 0.04$
200k CNN	$5.93 \pm 0.24$	$169 \pm 5.8$

## Training time

Model	Accuracy	Training time
150k S2CNN	97.64%	15h
5M CNN	97.49%	26h

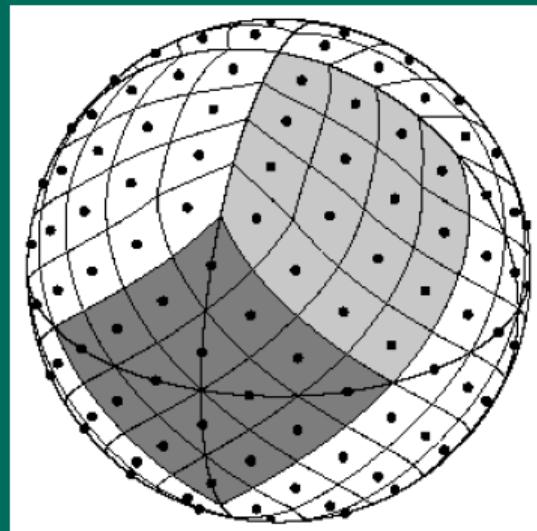
# Paper II: Conclusions and takeaways

Whether equivariance or data augmentation is better for an application depends on several factors:

- Which symmetries exists in your problem?
- Is the symmetry “amenable” to your model?
- How much data do you have access to?
- Is training time important?
- Is throughput / latency important?

# Future work

- Develop mathematical formalism for equivariant transformers
- Combining local and global symmetries
- Effects of sampling data with the HEALPix grid (Gorski et al. 2005)



Thank you

Thank you for listening!

# References I



‘Homogeneous Vector Bundles and  $\mathbb{G}$ -Equivariant Convolutional Neural Networks’. 11th May 2021. arXiv: 2105.05400 [cs, math, stat].



‘Roto-Translation Covariant Convolutional Networks for Medical Image Analysis’. 11th June 2018. arXiv: 1804.03393 [cs, math].



‘Geometric Deep Learning: Grids, Groups, Graphs, Geodesics, and Gauges’. 27th Apr. 2021. arXiv: 2104.13478 [cs, stat].



‘A General Theory of Equivariant CNNs on Homogeneous Spaces’. 9th Jan. 2020. arXiv: 1811.02017 [cs, stat].



‘Spherical CNNs’. 25th Feb. 2018. arXiv: 1801.10130 [cs, stat].



‘Group Equivariant Convolutional Networks’. 3rd June 2016. arXiv: 1602.07576 [cs, stat].

## References II



‘Geometric Deep Learning and Equivariant Neural Networks’. 28th May 2021. arXiv: 2105.13926 [hep-th].



*Equivariance versus Augmentation for Spherical Images*. 12th July 2022. doi: 10.48550/arXiv.2202.03990. arXiv: 2202.03990 [cs].



‘HEALPix – a Framework for High Resolution Discretization, and Fast Analysis of Data Distributed on the Sphere’. In: *The Astrophysical Journal* 622.2 (Apr. 2005), pp. 759–771. ISSN: 0004-637X, 1538-4357. doi: 10.1086/427976. arXiv: astro-ph/0409513.



*Natural Operations in Differential Geometry*. 1993rd edition. Berlin ; New York: Springer, 4th Feb. 1993. 440 pp. ISBN: 978-3-540-56235-1.



‘General E(2)-Equivariant Steerable CNNs’. 19th Nov. 2019. arXiv: 1911.08251 [cs, eess].

# References III



*Coordinate Independent Convolutional Networks*

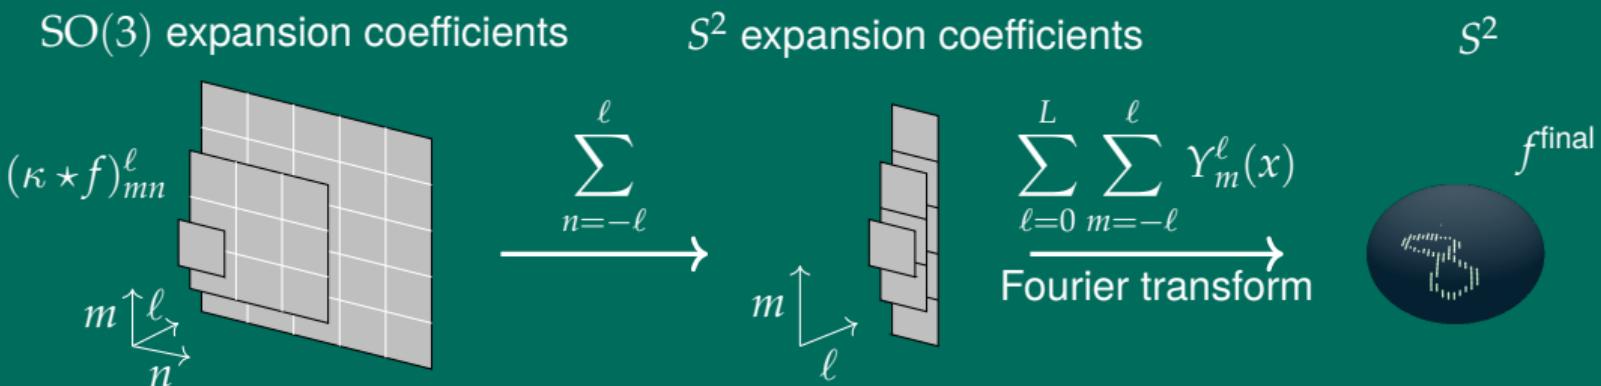
– *Isometry and Gauge Equivariant Convolutions on Riemannian Manifolds.*

10th June 2021. doi: 10.48550/arXiv.2106.06020. arXiv: 2106.06020 [cs, stat].

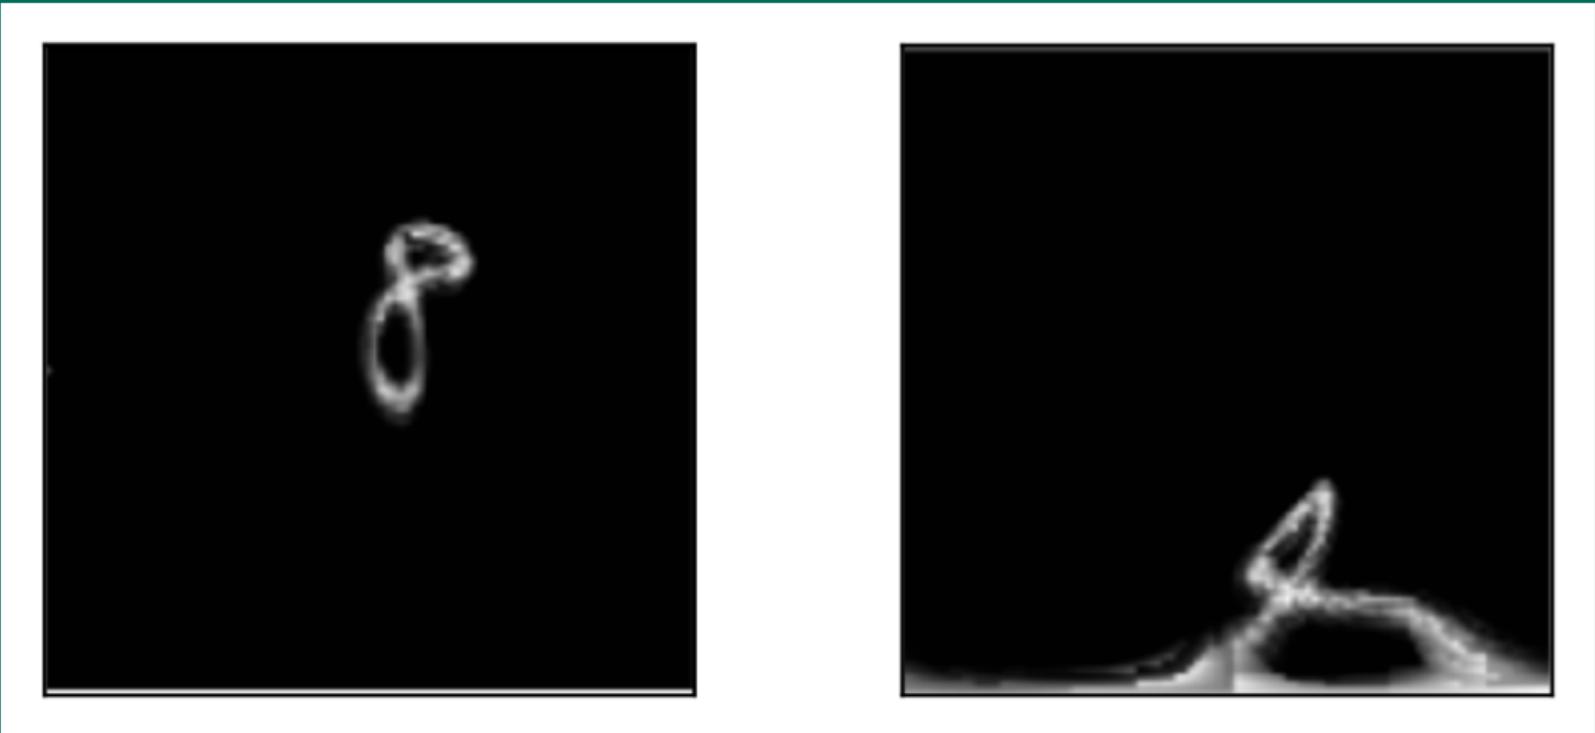


‘WoodScape: A Multi-Task, Multi-Camera Fisheye Dataset for Autonomous Driving’. 16th Aug. 2019. arXiv: 1905.01489 [cs, stat].

# Paper II: S2CNN extension



## Paper II: Rotated data points



# Mathematical formalism for supervised learning

Data: from  $\mathcal{X}$

Labels: from  $\mathcal{Y}$

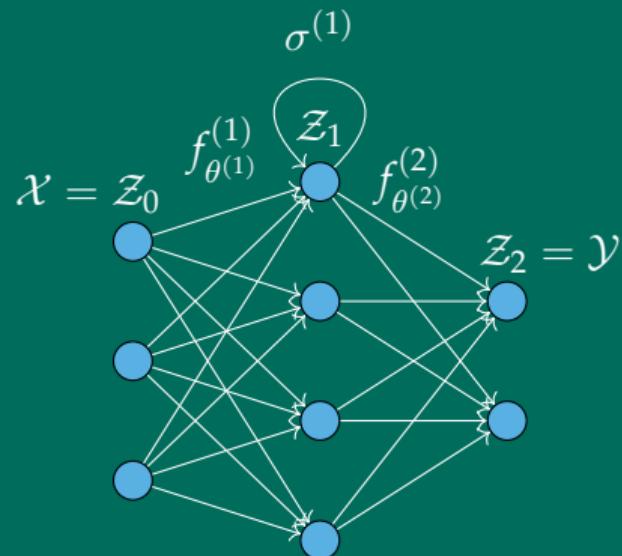
Layers:  $f_{\theta^{(l)}}^{(l)} : \mathcal{Z}_{l-1} \rightarrow \mathcal{Z}_l$

Non-linearity:  $\sigma^{(l)} : \mathcal{Z}_l \rightarrow \mathcal{Z}_l$

Model:  $\mathcal{N}_\theta : \mathcal{X} \rightarrow \mathcal{Y}, \quad \mathcal{N}_\theta = \bigcirc_{l=1}^N \sigma^{(l)} \circ f_{\theta^{(l)}}^{(l)}$

Dataset:  $D = \{(x_i, y_i)\}_{i=1\dots N} \subseteq \mathcal{X} \times \mathcal{Y}$

Loss function:  $\mathcal{L} : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$



Train the model by minimising the empirical risk through updating the parameters  $\theta$

$$\frac{1}{|D|} \sum_{(x_i, y_i) \in D} \mathcal{L}(\mathcal{N}_\theta(x_i), y_i)$$